1. Problem and Customer

Problem:

- ML engineers, researchers, and companies face high technical and decision-making overhead when fine-tuning or customizing LLMs.
- Existing platforms either focus on **frontend prototyping (e.g., v0.dev)** or **full-stack development (e.g., bolt.new)** but lack integrated research-backed tools for seamless **dataset discovery, model selection, and deployment**.
- Developers struggle to:
    - Find or generate suitable datasets.
    - Match datasets effectively with model architectures.
    - Rapidly iterate without rebuilding from scratch.

Customer:

- **Primary:** ML engineers, AI researchers, and startups building LLM-based products.
- **Secondary:** Enterprises investing in custom AI solutions that need internal prototyping capabilities.

---

2. Solution – How to Address the Problem

Agentic LLM Prototyping Platform

- Provides an **intelligent system for rapid prototyping and fine-tuning** of LLM/SLM models.
- Transforms abstract model concepts into **modular, deployable codebases** optimized for different compute environments.
- Integrates **deep research tooling** to help:
    - Identify optimal datasets.
    - Discover pretrained models.
    - Support architectural decision-making with multi-step reasoning.
- Built as **modular, composable tools** to solve high-friction tasks (starting with dataset preparation and model selection).

Key Components:

1. **Dataset Generation (v1 complete):** Task-specific, high-quality datasets with metadata and usage documentation.
2. **Research Agent (in progress):** Automates identification of best-fit datasets and pretrained models.

3. **Early Access Program:** Industry partners and developers test internal prototypes and provide feedback.

---

3. Justification – Customer Willingness to Pay

- **High Willingness to Pay:**
    - ML teams spend significant time and resources on dataset preparation and model fine-tuning. Reducing this friction provides immediate ROI.
    - Enterprises value faster iteration cycles and reduced engineering overhead.
- **Validation Evidence:**
    - Developers and researchers confirm dataset/model-matching is a **pain point worth paying for**.
    - Competing tools (e.g., Unsloth, Kiln AI, Gretel AI) show market demand, but lack the **research-driven modularity** this platform provides.
- **Strategic Timing:**
    - Recent advances in **small language models (SLMs)** enable efficient agent-driven workflows.
    - Market is still underexplored in **interactive research-as-a-service agents**, giving first-mover advantage.