Agentic LLM Prototyping Platform

One-Line Description

We are building a system similar to v0.dev, but focused on the training, fine-tuning, and deployment of LLM/SLM models, with deep research integration.

Expanded Overview

We are developing an intelligent system for rapid prototyping and fine-tuning of language models. The system supports transforming abstract model concepts into modular, deployable codebases optimized for different compute environments. This allows developers to test, iterate, and deploy without building from scratch.

The platform integrates deep research tools to assist in discovering relevant datasets, pretrained models, and architectural decisions.

Problem Statement

There is a high technical and decision-making overhead for ML engineers and companies trying to fine-tune or customize LLMs. Existing platforms either focus on full-stack development (e.g., bolt.new) or frontend prototyping (e.g., v0.dev) but lack tools for seamless research-backed model development.

Conversations with early-stage startup founders and portfolio investors from Nexus Ventures validated this gap to some extent. Our pivot came from developer feedback on our academic AI research platform—many struggled to find or generate suitable datasets and sought tooling to assist in dataset-model matching.

Short-Term Goals

We have divided the broader problem into modular, testable components:

- **Dataset Generation (Completed - v1):**

Deep-research-powered pipeline to create high-quality, task-specific datasets. Includes metadata and usage documentation.

- **Research Agent (In Progress):**

  A tool that identifies optimal datasets and pretrained models for a given task using multi-step reasoning and research evaluation.

- **Early Access Program:**

  We will soon onboard developers to test the platform and concurrently work with industry partners to build internal prototypes.

Strategic Rationale

- The broader challenge cannot be solved monolithically.
- We are chunking our solution space into well-defined, high-impact subproblems.
- Each tool is built with minimal dependencies and intended for composable use in research workflows.
- We prioritize solving high-friction tasks first (e.g., dataset preparation and model selection).

Timing Justification

- Recent progress in reasoning-capable small language models (SLMs) enables construction of non-trivial efficient and optimized agents.
- Research-as-a-service tooling has traditionally focused on static reports or dashboards; building interactive agents remain underexplored.

Similar products

- **Unsloth Notebooks:** Efficient GPU usage for fine-tuning tasks. Strong documentation. Focused on internal ops, not user modularity and finetuned model quality.

- **Kiln AI:** Provides low-code pipelines, but limited support for research-backed guidance.
- **DeepEval / Confident AI:** Strong in evaluation tooling but lacks system-level modularity for end-to-end prototyping.
- Mostly AI/ Gretel AI: can be a direct competitor.